


A crack team of love scientists from OkEros (a hot new dating site) have devised a way to represent dating profiles as rectangles on a two-dimensional plane.

They need help writing an algorithm to find the intersection of two users' love rectangles. They suspect finding that intersection is the key to a matching algorithm so *powerful* it will cause an immediate acquisition by Google or Facebook or Obama or something.



Write a function to find the rectangular intersection of two given love rectangles.

As with the example above, love rectangles are always "straight" and never "diagonal." More rigorously: each side is parallel with either the x-axis or the y-axis.

They are defined as hashes ↴

A **hash table** (also called a **hash**, **hash map**, **map**, **unordered map** or **dictionary**) is a data structure that pairs keys to values.

```
light_bulb_to_hours_of_light = {  
  'incandescent' => 1200,  
  'compact fluorescent' => 10000,  
  'LED' => 50000  
}
```

Ruby ▼

Hash tables:

- take *on average* **$O(1)$ time for insertions and lookups**
- are **unordered** (the keys are not guaranteed to stay in the same order)

- can use **many types of objects as keys** (commonly strings)

Hash tables can be thought of as arrays, if you think of array indices as keys!

In fact, hash tables are *built on* arrays. So if you ever want to use a hash table but know your keys will be sequential integers (like `1..100`), you can probably save time and space by just using an array instead.

Note: hash tables have an **average case** insertion and lookup cost of $O(1)$. In industry, we often confuse the average-case cost with *worst case* cost, but they're not really the same. Because of hash collisions and rebalancing, a hash table insertion or lookup can cost as much as $O(n)$ time in the worst case. But usually in industry we assume hashing and resizing algorithms are clever enough that collisions are rare and cheap.

like this:

```
my_rectangle = {  
  
  # coordinates of bottom-left corner  
  'left_x' => 1,  
  'bottom_y' => 5,  
  
  # width and height  
  'width' => 10,  
  'height' => 4,  
  
}
```

Ruby ▾

Your output rectangle should use this format as well.

Gotchas

What if there is *no* intersection? Does your function do something reasonable in that case?

What if one rectangle is entirely contained in the other? Does your function do something reasonable in that case?

What if the rectangles don't really intersect but share an edge? Does your function do something reasonable in that case?

Do some parts of your function seem very similar? Can they be refactored so you repeat yourself less?

Breakdown

Let's break this problem into subproblems. How can we divide this problem into smaller parts?

We could look at the two rectangles' "horizontal overlap" or "x overlap" separately from their "vertical overlap" or "y overlap."

Lets start with a helper function `find_x_overlap()`.

Need help finding the x overlap?

Since we're only working with the x dimension, we can treat the two rectangles' widths as ranges on a 1-dimensional number line.

What are the possible cases for how these ranges might overlap or not overlap? Draw out some examples!

There are four relevant cases:

1) The ranges partially overlap:



2) One range is completely contained in the other:



3) The ranges don't overlap:



4) The ranges "touch" at a single point:



Let's start with the first 2 cases. **How do we compute the overlapping range?**

One of our ranges starts "further to the right" than the other. We don't know ahead of time which one it is, but we can check the starting points of each range to see which one has the `highest_start_point`. **That `highest_start_point` is always the left-hand side of the overlap**, if there is one.

Not convinced? Draw some examples!

Similarly, **the right-hand side of our overlap is always the `lowest_end_point`**. That *may or may not* be the end point of the same input range that had the `highest_start_point`—compare cases (1) and (2).

This gives us our x overlap! So we can handle cases (1) and (2). **How do we know when there is *no* overlap?**

If `highest_start_point > lowest_end_point`, the two rectangles do not overlap.

But be careful—is it just *greater than* or is it *greater than or equal to*?

It depends how we want to handle case (4) above.

If we use *greater than*, we treat case (4) as an overlap. This means we could end up returning a rectangle with *zero width*, which ... may or may not be what we're looking for. You could make an argument either way.

Let's say a rectangle with zero width (or zero height) isn't a rectangle at all, so we should treat that case as "no intersection."

Can you finish `find_x_overlap()`?

Here's one way to do it:

```
def find_x_overlap(x1, width1, x2, width2)

  # find the highest ("rightmost") start point and lowest ("leftmost") end point
  highest_start_point = [x1, x2].max
  lowest_end_point = [x1 + width1, x2 + width2].min

  # return nil overlap if there is no overlap
  return [nil, nil] if highest_start_point >= lowest_end_point

  # compute the overlap width
  overlap_width = lowest_end_point - highest_start_point

  return [highest_start_point, overlap_width]
end
```

Ruby ▾

How can we adapt this for the rectangles' *ys* and *heights*?

Can we just make one `find_range_overlap()` function that can handle x overlap and y overlap?

Yes! We simply use more general parameter names:

```
def find_range_overlap(point1, length1, point2, length2)

  # find the highest start point and lowest end point.
  # the highest ("rightmost" or "upmost") start point is
  # the start point of the overlap.
  # the lowest end point is the end point of the overlap.
  highest_start_point = [point1, point2].max
  lowest_end_point = [point1 + length1, point2 + length2].min

  # return nil overlap if there is no overlap
  if highest_start_point >= lowest_end_point
    return [nil, nil]
  end

  # compute the overlap length
  overlap_length = lowest_end_point - highest_start_point

  return [highest_start_point, overlap_length]
end
```

We've solved our subproblem of finding the x and y overlaps! **Now we just need to put the results together.**

Solution

We divide the problem into two halves:

1. The intersection along the x-axis
2. The intersection along the y-axis

Both problems are basically the same as finding the intersection of two "ranges" on a 1-dimensional number line.

So we write a helper function `find_range_overlap()` that can be used to find both the x overlap and the y overlap, and we use it to build the rectangular overlap:

```
def find_range_overlap(point1, length1, point2, length2)

  # find the highest start point and lowest end point.
  # the highest ("rightmost" or "upmost") start point is
  # the start point of the overlap.
  # the lowest end point is the end point of the overlap.
  highest_start_point = [point1, point2].max
  lowest_end_point = [point1 + length1, point2 + length2].min

  # return nil overlap if there is no overlap
  if highest_start_point >= lowest_end_point
    return [nil, nil]
  end

  # compute the overlap length
  overlap_length = lowest_end_point - highest_start_point

  return [highest_start_point, overlap_length]
end

def find_rectangular_overlap(rect1, rect2)

  # get the x and y overlap points and lengths
  x_overlap_point, overlap_width = find_range_overlap(\
    rect1['left_x'], rect1['width'], rect2['left_x'], rect2['width'])
  y_overlap_point, overlap_height = find_range_overlap(\
    rect1['bottom_y'], rect1['height'], rect2['bottom_y'], rect2['height'])

  # return nil rectangle if there is no overlap
  if !overlap_width || !overlap_height
    return {
      'left_x' => nil,
      'bottom_y' => nil,
      'width' => nil,
      'height' => nil,
    }
  end

  return {
    'left_x' => x_overlap_point,
    'bottom_y' => y_overlap_point,
    'width' => overlap_width,
    'height' => overlap_height,
  }
end
```

Complexity

$O(1)$ time and $O(1)$ space.

Bonus

What if we had an array of rectangles and wanted to find *all* the rectangular overlaps between all possible pairs of two rectangles within the array? Note that we'd be returning *an array of rectangles*.

What if we had an array of rectangles and wanted to find the overlap between *all* of them, if there was one? Note that we'd be returning *a single rectangle*.

What We Learned

This is an interesting one because the hard part isn't the time or space optimization—it's getting something that *works* and is *readable*.

For problems like this, I often see candidates who can describe the strategy at a high level but trip over themselves when they get into the details.

Don't let it happen to you. To keep your thoughts clear and avoid bugs, take time to:

1. Think up and draw out all the possible cases. Like we did with the ways ranges can overlap.
2. Use very specific and descriptive variable names.

Want more coding interview help?

Check out **interviewcake.com** for more advice, guides, and practice questions.